

```
In [6]: import pandas as pd
import numpy as np

# Load the data (adjust path if needed)
data = pd.read_csv('../garments_worker_productivity.csv')
print(data.head())
print(data.columns)
```

	date	quarter	department	day	team	targeted_productivity	\
0	1/1/2015	Quarter1	sweing	Thursday	8	0.80	
1	1/1/2015	Quarter1	finishing	Thursday	1	0.75	
2	1/1/2015	Quarter1	sweing	Thursday	11	0.80	
3	1/1/2015	Quarter1	sweing	Thursday	12	0.80	
4	1/1/2015	Quarter1	sweing	Thursday	6	0.80	

	smv	wip	over_time	incentive	idle_time	idle_men	\
0	26.16	1108.0	7080	98	0.0	0	
1	3.94	NaN	960	0	0.0	0	
2	11.41	968.0	3660	50	0.0	0	
3	11.41	968.0	3660	50	0.0	0	
4	25.90	1170.0	1920	50	0.0	0	

	no_of_style_change	no_of_workers	actual_productivity
0	0	59.0	0.940725
1	0	8.0	0.886500
2	0	30.5	0.800570
3	0	30.5	0.800570
4	0	56.0	0.800382

```
Index(['date', 'quarter', 'department', 'day', 'team', 'targeted_productivity',
      'smv', 'wip', 'over_time', 'incentive', 'idle_time', 'idle_men',
      'no_of_style_change', 'no_of_workers', 'actual_productivity'],
      dtype='object')
```

```
In [7]: # Check for missing values
print(data.isnull().sum())
print("\nShape before dropping columns with too many nulls:", data.shape)

# Drop columns with more than 20% missing values (if any)
data = data.dropna(axis=1, thresh=len(data)*0.8)
print("\nShape after dropping columns with too many nulls:", data.shape)

# Fill remaining missing values
for col in data.select_dtypes(include='O').columns:
    data[col] = data[col].fillna(data[col].mode()[0])
for col in data.select_dtypes(include=[np.number]).columns:
    data[col] = data[col].fillna(data[col].mean())

# Check again
print("\nMissing values after filling:")
print(data.isnull().sum())
```

```

date            0
quarter         0
department      0
day            0
team           0
targeted_productivity  0
smv            0
wip            506
over_time       0
incentive       0
idle_time       0
idle_men        0
no_of_style_change  0
no_of_workers   0
actual_productivity  0
dtype: int64

```

Shape before dropping columns with too many nulls: (1197, 15)

Shape after dropping columns with too many nulls: (1197, 14)

Missing values after filling:

```

date            0
quarter         0
department      0
day            0
team           0
targeted_productivity  0
smv            0
over_time       0
incentive       0
idle_time       0
idle_men        0
no_of_style_change  0
no_of_workers   0
actual_productivity  0
dtype: int64

```

```

In [8]: # Convert 'date' to datetime and extract 'month'
data['date'] = pd.to_datetime(data['date'])
data['month'] = data['date'].dt.month

# Drop the original 'date' column
data = data.drop('date', axis=1)

# Show the first few rows and columns
print(data.head())
print("\nColumns after date processing:", data.columns.tolist())

```

	quarter	department	day	team	targeted_productivity	smv	\
0	Quarter1	sweing	Thursday	8	0.80	26.16	
1	Quarter1	finishing	Thursday	1	0.75	3.94	
2	Quarter1	sweing	Thursday	11	0.80	11.41	
3	Quarter1	sweing	Thursday	12	0.80	11.41	
4	Quarter1	sweing	Thursday	6	0.80	25.90	

	over_time	incentive	idle_time	idle_men	no_of_style_change	\
0	7080	98	0.0	0	0	
1	960	0	0.0	0	0	
2	3660	50	0.0	0	0	
3	3660	50	0.0	0	0	
4	1920	50	0.0	0	0	

	no_of_workers	actual_productivity	month
0	59.0	0.940725	1
1	8.0	0.886500	1
2	30.5	0.800570	1
3	30.5	0.800570	1
4	56.0	0.800382	1

Columns after date processing: ['quarter', 'department', 'day', 'team', 'targeted_productivity', 'smv', 'over_time', 'incentive', 'idle_time', 'idle_men', 'no_of_style_change', 'no_of_workers', 'actual_productivity', 'month']

```
In [9]: from sklearn.preprocessing import LabelEncoder

# Encode 'quarter', 'department', and 'day'
for col in ['quarter', 'department', 'day']:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])

print(data.head())
print("\nData types after encoding:\n", data.dtypes)
```

	quarter	department	day	team	targeted_productivity	smv	over_time	\
0	0	2	3	8	0.80	26.16	7080	
1	0	1	3	1	0.75	3.94	960	
2	0	2	3	11	0.80	11.41	3660	
3	0	2	3	12	0.80	11.41	3660	
4	0	2	3	6	0.80	25.90	1920	

	incentive	idle_time	idle_men	no_of_style_change	no_of_workers	\
0	98	0.0	0	0	59.0	
1	0	0.0	0	0	8.0	
2	50	0.0	0	0	30.5	
3	50	0.0	0	0	30.5	
4	50	0.0	0	0	56.0	

	actual_productivity	month
0	0.940725	1
1	0.886500	1
2	0.800570	1
3	0.800570	1
4	0.800382	1

Data types after encoding:

```

quarter          int64
department       int64
day             int64
team            int64
targeted_productivity float64
smv             float64
over_time       int64
incentive       int64
idle_time       float64
idle_men        int64
no_of_style_change int64
no_of_workers   float64
actual_productivity float64
month           int32
dtype: object

```

```

In [10]: from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = data.drop(['actual_productivity'], axis=1)
y = data['actual_productivity']

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)

print("Feature columns:", X.columns.tolist())
print("Number of features:", len(X.columns))
print("Train shape:", X_train.shape)
print("Test shape:", X_test.shape)

```

```

Feature columns: ['quarter', 'department', 'day', 'team', 'targeted_productivity', 'smv', 'over_time', 'incentive', 'idle_time', 'idle_men', 'no_of_style_change', 'no_of_workers', 'month']
Number of features: 13
Train shape: (957, 13)
Test shape: (240, 13)

```

```

In [11]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Linear Regression

```

```

lr = LinearRegression()
lr.fit(X_train, y_train)
pred_lr = lr.predict(X_test)

# Random Forest
rf = RandomForestRegressor(random_state=42)
rf.fit(X_train, y_train)
pred_rf = rf.predict(X_test)

# XGBoost
xgb = XGBRegressor(random_state=42)
xgb.fit(X_train, y_train)
pred_xgb = xgb.predict(X_test)

# Model Comparison
print('Model Comparison:')
models = ['Linear Regression', 'Random Forest', 'XGBoost']
preds = [pred_lr, pred_rf, pred_xgb]
for name, pred in zip(models, preds):
    print(f'--- {name} ---')
    print('MAE:', mean_absolute_error(y_test, pred))
    print('MSE:', mean_squared_error(y_test, pred))
    print('R2 Score:', r2_score(y_test, pred))
    print()

```

Model Comparison:

--- Linear Regression ---

MAE: 0.10750872834094855

MSE: 0.02160847942846644

R2 Score: 0.18619690531408462

--- Random Forest ---

MAE: 0.0668503848365833

MSE: 0.011728750658163357

R2 Score: 0.5582801828323646

--- XGBoost ---

MAE: 0.07273557406922082

MSE: 0.01505318533673072

R2 Score: 0.4330777020906966

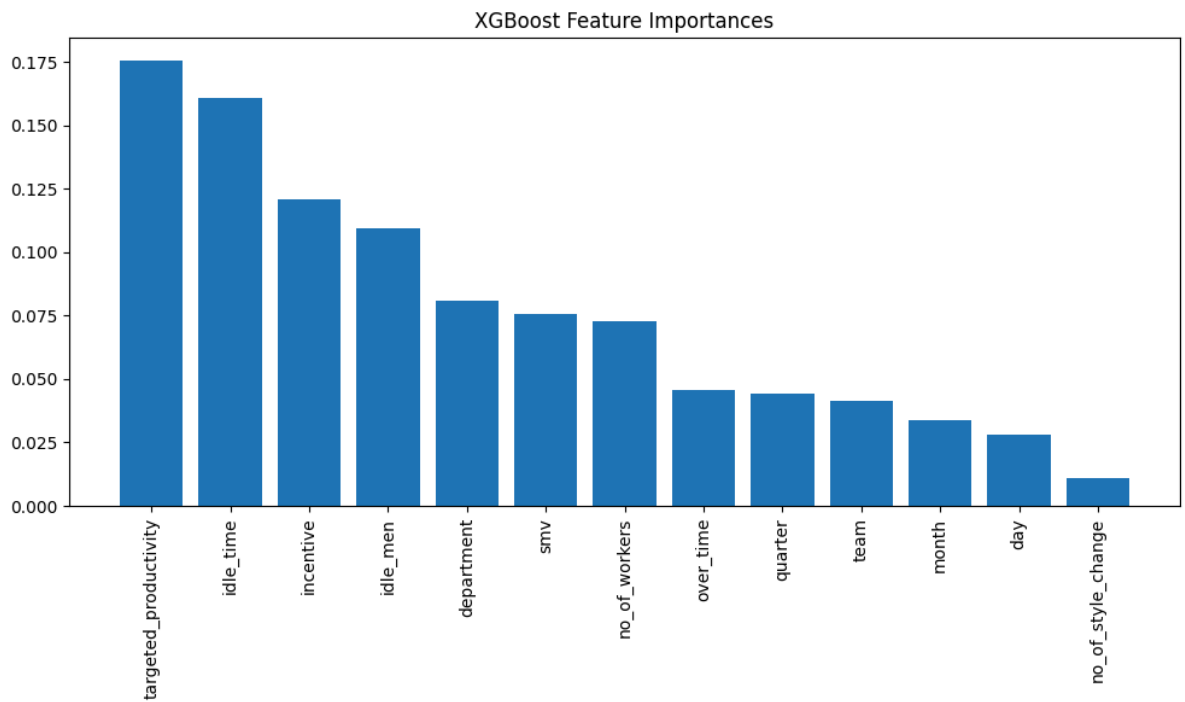
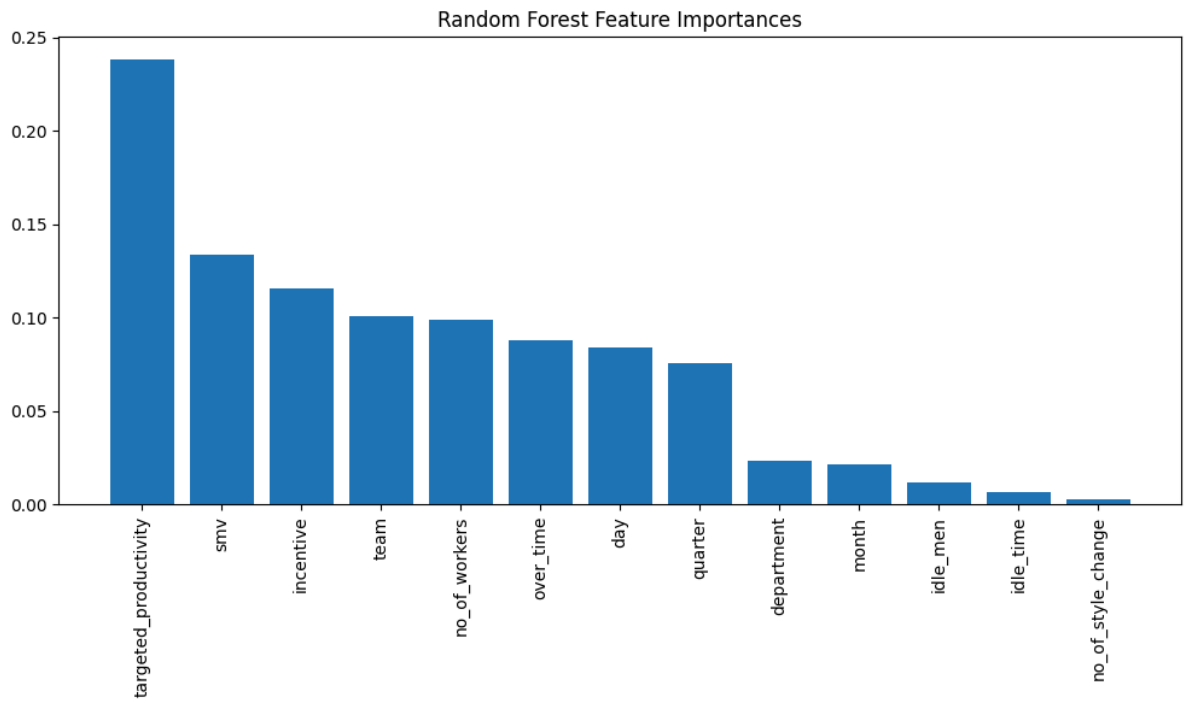
```

In [12]: import matplotlib.pyplot as plt
import numpy as np

# Random Forest Feature Importance
importances_rf = rf.feature_importances_
indices_rf = np.argsort(importances_rf)[::-1]
plt.figure(figsize=(10,6))
plt.title('Random Forest Feature Importances')
plt.bar(range(X.shape[1]), importances_rf[indices_rf], align='center')
plt.xticks(range(X.shape[1]), X.columns[indices_rf], rotation=90)
plt.tight_layout()
plt.show()

# XGBoost Feature Importance
importances_xgb = xgb.feature_importances_
indices_xgb = np.argsort(importances_xgb)[::-1]
plt.figure(figsize=(10,6))
plt.title('XGBoost Feature Importances')
plt.bar(range(X.shape[1]), importances_xgb[indices_xgb], align='center')
plt.xticks(range(X.shape[1]), X.columns[indices_xgb], rotation=90)
plt.tight_layout()
plt.show()

```



```
In [13]: import pickle

# Save the Random Forest model
with open('./gwp.pkl', 'wb') as f:
    pickle.dump(rf, f)

print("Random Forest model saved as gwp.pkl")
```

Random Forest model saved as gwp.pkl